

Docker lab

Lab-1

Install docker on your linux virtual system

<https://docs.docker.com/engine/install/fedora/>

Start hello-world container to test your installation

```
[root@server ~]# docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
1b930d010525: Pull complete
Digest: sha256:c3b4ada4687bbaa170745b3e4dd8ac3f194ca95b2d0518b417fb47e5879d9b5f
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker!

This message shows that your installation appears to be working correctly.

[...]

Inspect images and container with docker cli

Get container active:

```
[root@server ~]# docker ps
```

Get all the container

```
[root@server ~]# docker ps -a
```

Get container images downloaded locally

```
[root@server ~]# docker images
```

Search images on docker hub

```
[root@server ~]# docker search hello-world
```

Delete hello-world container

```
[root@server ~]# docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	PORTS
CREATED	STATUS		
NAMES			
9fdbdd35f5a5	hello-world	"/hello"	7
minutes ago	Exited (0) 7 minutes ago		
relaxed_shtern			

Delete hello-world container

```
[root@server ~]# docker rm 9fdbdd35f5a5
9fdbdd35f5a5
```

Delete hello-world image

```
[root@server ~]# docker rmi hello-world
```

Lab-2

Let's install NGINX with the following command:

```
[root@server ~]# docker run --detach --name web_test nginx:latest
Unable to find image 'nginx:latest' locally
latest: Pulling from library/nginx
8d691f585fa8: Pull complete
5b07f4e08ad0: Pull complete
abc291867bca: Pull complete
Digest: sha256:922c815aa4df050d4df476e92daed4231f466acc8ee90e0e774951b0fd7195a4
Status: Downloaded newer image for nginx:latest
e6cdf5e34c0900c7db3002424bba8d88dc1c7c50c79893820a7bff0ee9c3934a
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
e6cdf5e34c09	nginx:latest	"nginx -g 'daemon of..."	7 seconds ago
Up 6 seconds	80/tcp	web_test	

Stop docker container web_test

```
[root@server ~]# docker stop web_test
web_test
```

Start again a new nginx container binding port 80 on local system

```
[root@server ~]# docker run --detach -p 80:80 --name web nginx:latest
f8baa546333981995b13ce4d6be5478030ecbbb2fc266d4a9b12ce8199fa1077
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
f8baa5463339	nginx:latest	"nginx -g 'daemon of..."	44 seconds ago
Up 43 seconds	0.0.0.0:80->80/tcp	web	

Try to open a http connection from your local browser on forwarded port 8080 (virtualbox system)



Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

Inspect container log to see the connection log entry

```
[root@server ~]# docker logs web
```

```
192.168.3.2 - - [15/Nov/2019:09:41:29 +0000] "GET / HTTP/1.1" 200 612 "-"
"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_14_6) AppleWebKit/605.1.15 (KHTML, like
Gecko) Version/13.0.3 Safari/605.1.15" "-"
```

Stop web container

```
[root@server ~]# docker stop web
```

web

```
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	

LAB-3

Start another container named web --name web:

```
[root@server ~]# docker run --detach --name web nginx:latest
```

docker: Error response from daemon: Conflict. The container name "/web" is already in use by container

"f8baa546333981995b13ce4d6be5478030ecbbb2fc266d4a9b12ce8199fa1077". You have to remove (or rename) that container to be able to reuse that name.

See 'docker run --help'.

Solution 1: rename

```
[root@server ~]# docker rename web web_old
```

```
[root@server ~]# docker run --detach --name web nginx:latest
```

5277054d7d7a18c37308feb6cb7e6d629cccaf55de898c4e683c91975f37b40e

Solution 2: stop & remove

```
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
5277054d7d7a	nginx:latest	"nginx -g 'daemon of..."	53 seconds ago
Up 52 seconds	80/tcp	web	

```
[root@server ~]# docker stop web
```

web

```
[root@server ~]# docker rm web
```

web

```
[root@server ~]# docker run --detach --name web nginx:latest
```

12d981cfb7288f127bc15625bfea420cac06c2f3c649e2b2c35312e29e4bae0c

```
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
12d981cfb728	nginx:latest	"nginx -g 'daemon of..."	2 seconds ago
Up 2 seconds	80/tcp	web	

LAB-5

Pull images from Docker Registry

```
[root@server ~]# docker pull httpd:latest
latest: Pulling from library/httpd
8d691f585fa8: Pull complete
8eb779d8bd44: Pull complete
574add29ec5c: Pull complete
30d7fa9ec230: Pull complete
ede292f2b031: Pull complete
Digest:
sha256:35fcab73dc9ae55db5c4ac33f5e0c7e76b7735aaddb628366bab04db6f8
ae96e
Status: Downloaded newer image for httpd:latest
docker.io/library/httpd:latest
```

Verify:

```
[root@server ~]# docker images
```

REPOSITORY	TAG	IMAGE ID	
httpd	latest	d3017f59d5e2	2
weeks ago	165MB		

Export image to a local file

```
[root@server ~]# docker save -o httpd_latest.tar httpd:latest
[root@server ~]# ls -la httpd_latest.tar
-rw-----. 1 root root 170482176 20 nov 05.25 httpd_latest.tar
```

Remove httpd image

```
[root@server ~]# docker rmi httpd
Untagged: httpd:latest
Untagged:
httpd@sha256:35fcab73dc9ae55db5c4ac33f5e0c7e76b7735aaddb628366bab0
4db6f8ae96e
Deleted:
sha256:d3017f59d5e25daba517ac35eaf4b862dce70d2af5f27bf40bef5f936c8
b2e1f
Deleted:
sha256:c015bdd664253fc2ccdf3a425ba085e94b99ce801d6c9c5219ffaad279
362b1
Deleted:
sha256:c79505e64684e42a92353a6e3430969d8a801f327d24bdde11bd99d41cb
ef2a0
Deleted:
sha256:87158971c3545200bc870118cfa31bf5470204eea10da0b79531388b5f9
1cea2
```

Deleted:
sha256:0105db3a8b98aea80771956b067b64a26bd630718141bacd6feafdc5b5c0caee

Deleted:
sha256:b67d19e65ef653823ed62a5835399c610a40e8205c16f839c5cc567954fcf594

```
[root@server ~]# docker images
```

REPOSITORY	TAG	IMAGE ID
CREATED	SIZE	

Load local images from file

```
[root@server ~]# docker load -i httpd_latest.tar
```

b67d19e65ef6: Loading layer

[=====>]

72.5MB/72.5MB

71b4839cc8bf: Loading layer

[=====>]

2.56kB/2.56kB

e60e854dd58d: Loading layer

[=====>]

36.7MB/36.7MB

f69951420260: Loading layer

[=====>]

61.24MB/61.24MB

108b118a67d7: Loading layer

[=====>]

3.584kB/3.584kB

Loaded image: httpd:latest

```
[root@server ~]# docker images
```

REPOSITORY	TAG	IMAGE ID	
CREATED	SIZE		
httpd	latest	d3017f59d5e2	2
weeks ago	165MB		

Lab-6

Playing with Busybox

In this lab, we are going to run a [Busybox](#) container on our system and get a taste of the docker run command.

```
[root@server ~]# docker pull busybox
Using default tag: latest
latest: Pulling from library/busybox
0f8c40e1270f: Pull complete
Digest:
sha256:1303dbf110c57f3edf68d9f5a16c082ec06c4cf7604831669faf2c71226
0b5a0
Status: Downloaded newer image for busybox:latest
docker.io/library/busybox:latest
[root@server ~]# docker images
```

REPOSITORY	TAG	IMAGE ID	
CREATED	SIZE		
busybox	latest	020584afccce	2
weeks ago	1.22MB		

The pull command fetches the busybox image from the Docker registry and saves it to our system. You can use the docker images command to see a list of all images on your system.

```
[root@server ~]# docker images
```

REPOSITORY	TAG	IMAGE ID	
CREATED	SIZE		
busybox	latest	020584afccce	2
weeks ago	1.22MB		

Run a Docker container based on this image:

```
[root@server ~]# docker run busybox
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	
CREATED	STATUS	PORTS	NAMES

nothing happened! Is that a bug?... no!. Behind the scenes, a lot of stuff happened. When you call run, the Docker client finds the image (busybox in this case), loads up the container and then runs a command in that container. When we run docker run busybox, we didn't provide a command, so the container booted up, ran an empty command and then exited.

run this command:

```
[root@server ~]# docker run busybox echo "hello from busybox"
hello from busybox
```

In this case, the Docker client ran the **echo** command in our busybox container and then exited it. If you've noticed, all of that happened pretty quickly. Imagine booting up a virtual machine, running a command and then killing it.

```
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND
CREATED	STATUS	PORTS
NAMES		

```
[root@server ~]# docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	
CREATED	STATUS	PORTS	
NAMES			
7bdbcb110894	busybox	"echo 'hello from bu..."	2
minutes ago	Exited (0)	2 minutes ago	
stupefied_bell			
ec46a2d08d67	busybox	"sh"	4
hours ago	Exited (0)	4 hours ago	
naughty_nash			

To run more than just one command in a container. Let's try that now:

```
[root@server ~]# docker run -it busybox sh
/ # ls
bin    dev    etc    home  proc  root  sys   tmp   usr   var
/ # uptime
 14:50:23 up   4:32,   0 users,   load average: 0.00, 0.01, 0.04
/ #
```

Running the run command with the -it flags attaches us to an interactive tty in the container. Now we can run as many commands in the container as we want. **Take some time to run some commands.**

Clean up all the container exited from start to now

```
[root@server ~]# docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	
CREATED	STATUS	PORTS	
NAMES			
73fab3ef1564	busybox	"sh"	
16 minutes ago	Exited (130)	58 seconds ago	
priceless_lovelace			
7bdbcb110894	busybox	"echo 'hello from bu..."	
20 minutes ago	Exited (0)	20 minutes ago	
stupefied_bell			


```

0ac16a224243      hello-world      "/hello"      4
hours ago      Exited (0) 4 hours ago
priceless_cerf
ec46a2d08d67      busybox      "sh"      4
hours ago      Exited (0) 4 hours ago
naughty_nash
ef86117d695d      d3017f59d5e2      "httpd-foreground"      4
hours ago      Exited (0) 4 hours ago
blissful_austin
3fb98025290b      d3017f59d5e2      "httpd-foreground"      4
hours ago      Exited (0) 4 hours ago
unruffled_mclean
[root@server ~]# docker ps -a -q -f status=exited
73fab3ef1564
7bdbcb110894
0ac16a224243
ec46a2d08d67
ef86117d695d
3fb98025290b
[root@server ~]# docker rm $(docker ps -a -q -f status=exited)
73fab3ef1564
7bdbcb110894
0ac16a224243
ec46a2d08d67
ef86117d695d
3fb98025290b
[root@server ~]# docker ps -a
CONTAINER ID        IMAGE               COMMAND
CREATED            STATUS             PORTS             NAMES

```

Terminology

Some terminology used in this exercise on the Docker ecosystem.

Images - The blueprints of our application which form the basis of containers. In the demo above, we used the `docker pull` command to download the busybox image.

Containers - Created from Docker images and run the actual application. We create a container using `docker run` which we did using the busybox image that we downloaded. A list of running containers can be seen using the `docker ps` command.

Docker Daemon - The background service running on the host that manages building, running and distributing Docker containers. The daemon is the process that runs in the operating system to which clients talk to.

Docker Client - The command line tool that allows the user to interact with the daemon. More generally, there can be other forms of clients too - such as Kitematic which provide a GUI to the users.

Docker Hub - A registry of Docker images. You can think of the registry as a directory of all available Docker images. If required, one can host their own Docker registries and can use them for pulling images.

LAB-7


Run a static website in a container

The image that you are going to use is a single-page website available on the Docker Store as [dockersamples/static-site](https://store.docker.com/images/docker-samples-static-site). You can download and run the image directly in one go using `docker run` as follows.

```
[root@server ~]# docker run --name static-site-2 -e AUTHOR="Your Name" -d -p 80:80 dockersamples/static-site
fe79037a930c91b2e5b548cf369c2ee92cf6d359f0a545326cd927433464e016
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND
CREATED	STATUS	PORTS
NAMES		
fe79037a930c	dockersamples/static-site	"/bin/sh -c 'cd /usr..."
3 seconds ago	Up 2 seconds	
0.0.0.0:80->80/tcp, 443/tcp	static-site-2	

Try to open a http connection from your local browser on forwarded port 8080 (virtualbox system)



Hello Luca Cavatorta!

This is being served from a **docker** container running Nginx.

Now, let's launch a container in detached mode as shown below:

```
[root@server ~]# docker run --name static-site -e AUTHOR="Luca Cavatorta" -d -P dockersamples/static-site
4f1d765fa9a6eec01b62f455c59199666b804cff937c553a104f180abb4a36a0
```

e61d12292d69556eabe2a44c16cbd54486b2527e2ce4f95438e504afb7b02810

In the above command:

- -d will create a container with the process detached from our terminal
- -P will publish all the exposed container ports to **random** ports on the Docker host
- -p will expose your container port selected ports on the Docker host
- **-e is how you pass environment variables to the container**
- --name allows you to specify a container name
- AUTHOR is the environment variable name and Your Name is the value that you can pass

Now you can see the ports by running the docker port command.

```
[root@server ~]# docker port static-site
443/tcp -> 0.0.0.0:32768
80/tcp -> 0.0.0.0:32769
```

If you want create a temporary bind on virtualbox to port 32769 and 32768 and try to open a local browser on forward port configured.

Clean up all

docker ps

docker stop

docker ps -a

docker rm

LAB-8

Volumes:

Bind mount volume

Create a new directory /localvolume

Add some file to /localvolume directory

for example:

```
touch /localvolume/file1
```

```
touch /localvolume/file2
```

Start a fedora container bind to it your /localvolume directory

```
[root@server ~]# docker run -it --name myfedora -v
/localvolume:/usr/localvolume:ro fedora bash
[root@fb6a2fd1de6d /]# cd /usr/localvolume
[root@fb6a2fd1de6d localvolume]# ls -la
file1
file2
```

Try to create a directory or a file under /usr/localvolume

Stop the container

Docker-managed volume and shared volumes

Start a fedora container named myfedora2 with a new Docker volume

```
docker run -dit --name myfedora2 -v /localvolume fedora bash
Inspect create volume
```

Start a second fedora container named myfedora3 with shared volume from myfedora2 container `docker run -it -d --name myfedora3 --volumes-from myfedora2 fedora`

```
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	
30efc2065fa5	fedora	"bash"	3
seconds ago	Up 2 seconds		
myfedora3			

6931f31807f4	fedora	"bash"	2
minutes ago	Up 2 minutes		
myfedora2			

Create some file on myfedora2 container under /localvolume

```
[root@server ~]# docker exec -it myfedora2 bash
[root@6931f31807f4 /]# cd /localvolume/
[root@6931f31807f4 localvolume]# ll
total 0
[root@6931f31807f4 localvolume]# touch file1
[root@6931f31807f4 localvolume]# ls -la
total 0
-rw-r--r--. 1 root root 0 Nov 21 10:12 file1
[root@6931f31807f4 localvolume]# exit
exit
```

Verify on myfedora3 /localvolume directory file

```
[root@server ~]# docker exec -it myfedora3 bash
[root@30efc2065fa5 /]# cd /localvolume/
[root@30efc2065fa5 localvolume]# ls -la
total 0
drwxr-xr-x. 2 root root 19 Nov 21 10:12 .
drwxr-xr-x. 1 root root 25 Nov 21 10:11 ..
-rw-r--r--. 1 root root 0 Nov 21 10:12 file1
```

Stop myfedora3 myfedora2 container and delete created docker volume

```
[root@server ~]# docker volume ls
DRIVER          VOLUME NAME
local
395b9c327d9fcad5cbea9401eb10ef28f492b3e8b7e3b79d6a7ef28f5e73494b
[root@server ~]# docker volume rm
395b9c327d9fcad5cbea9401eb10ef28f492b3e8b7e3b79d6a7ef28f5e73494b
```

If are you sure that no other volumes present are used, to delete all non used volume you can use:

```
docker volume prune
WARNING! This will remove all local volumes not used by at least
one container.
Are you sure you want to continue? [y/N]
Deleted Volumes:
395b9c327d9fcad5cbea9401eb10ef28f492b3e8b7e3b79d6a7ef28f5e73494b
```

```
Total reclaimed space: 0B
[root@server ~]# docker volume ls
DRIVER          VOLUME NAME
```

LAB-9

Communication across links

Links allow containers to discover each other and securely transfer information about one container to another container. When you set up a link, you create a conduit between a source container and a recipient container. The recipient can then access select data about the source. To create a link, you use the `--link` flag. First, create a new container, this time one containing a database.

```
$ docker run -d --name db training/postgres
```

This creates a new container called **db** from the **training/postgres** image, which contains a **PostgreSQL** database.

Now, create a new **web** container and link it with your db container.

```
$ docker run -d -P --name web --link db:db training/webapp python
app.py
```

This links the new **web** container with the **db** container you created earlier. The `--link` flag takes the form:

`--link <name or id>:alias`

Next, inspect your linked containers with `docker inspect`:

```
$ docker inspect -f "{{ .HostConfig.Links }}" web

[/db:/web/db]
```

You can see that the **web** container is now linked to the **db** container **web/db**. Which allows it to access information about the db container.

```
$ docker exec -it web env
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=4b3bbafd5430
TERM=xterm
DB_PORT=tcp://172.17.0.3:5432
```

DB_PORT_5432_TCP=tcp://172.17.0.3:5432
DB_PORT_5432_TCP_ADDR=172.17.0.3
DB_PORT_5432_TCP_PORT=5432
DB_PORT_5432_TCP_PROTO=tcp
DB_NAME=/web/db
DB_ENV_PG_VERSION=9.3
HOME=/root

LAB-10

docker attach

Attach local standard input, output, and error streams to a running container

Use **docker attach** to attach your terminal's standard input, output, and error (or any combination of the three) to a running container using the container's ID or name. This allows you to view its ongoing output or to control it interactively, as though the commands were running directly in your terminal.

Note: The **attach** command will display the output of the **ENTRYPOINT/CMD** process. This can appear as if the attach command is hung when in fact the process may simply not be interacting with the terminal at that time.

You can attach to the same contained process multiple times simultaneously, from different sessions on the Docker host.

To **stop** a container, use **CTRL-c**. This key sequence sends **SIGKILL** to the container. If **--sig-proxy** is true (the default), **CTRL-c** sends a **SIGINT** to the container. If the container was run with **-i** and **-t**, you can detach from a container and leave it running using the **CTRL-p CTRL-q** key sequence.

Note: A process running as PID 1 inside a container is treated specially by Linux: it ignores any signal with the default action. So, the process will not terminate on SIGINT or SIGTERM unless it is coded to do so.

Use user-defined bridge networks

In this example, we again start two alpine containers, but attach them to a user-defined network called alpine-net which we have already created. These containers are not connected to the default bridge network at all. We then start a third alpine container which is connected to the bridge network but not connected to alpine-net, and a fourth alpine container which is connected to both networks.

Create the alpine-net network. You do not need the **--driver bridge** flag since it's the default, but this example shows how to specify it.

```
$ docker network create --driver bridge alpine-net
```

List Docker's networks:

```
$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
e9261a8c9a19	alpine-net	bridge	local

17e324f45964	bridge	bridge	local
6ed54d316334	host	host	local
7092879f2cc8	none	null	local

Inspect the alpine-net network. This shows you its IP address and the fact that no containers are connected to it:

```
$ docker network inspect alpine-net
```

```
[
  {
    "Name": "alpine-net",
    "Id":
    "e9261a8c9a19eabf2bf1488bf5f208b99b1608f330cff585c273d39481c9b0ec"
  ,
    "Created": "2017-09-25T21:38:12.620046142Z",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": {},
      "Config": [
        {
          "Subnet": "172.18.0.0/16",
          "Gateway": "172.18.0.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Containers": {},
    "Options": {},
    "Labels": {}
  }
]
```

Notice that this network's gateway is **172.18.0.1**, as opposed to the default bridge network, whose gateway is **172.17.0.1** (`docker network inspect bridge`). The exact IP address may be different on your system.

Create your four containers. Notice the `--network` flags. You can only connect to one network during the `docker run` command, so you need to use `docker network connect` afterward to connect `alpine4` to the bridge network as well.

```
$ docker run -dit --name alpine1 --network alpine-net alpine ash
```

```
$ docker run -dit --name alpine2 --network alpine-net alpine ash
```

```
$ docker run -dit --name alpine3 alpine ash
```

```
$ docker run -dit --name alpine4 --network alpine-net alpine ash
```

```
$ docker network connect bridge alpine4
```

Verify that all containers are running:

```
$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	NAMES
156849ccd902	alpine	"ash"	alpine4
seconds ago	Up 41 seconds		
fa1340b8d83e	alpine	"ash"	alpine3
seconds ago	Up 51 seconds		
a535d969081e	alpine	"ash"	alpine2
a minute ago	Up About a minute		
0a02c449a6e9	alpine	"ash"	alpine1
a minute ago	Up About a minute		

Inspect the bridge network and the alpine-net network again:

```
$ docker network inspect bridge
```

```
[
  {
    "Name": "bridge",
    "Id": "17e324f459648a9baaea32b248d3884da102dde19396c25b30ec800068ce6b10"
  },
  {
    "Created": "2017-06-22T20:27:43.826654485Z",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "172.17.0.0/16",
          "Gateway": "172.17.0.1"
        }
      ]
    }
  }
]
```

```

    ],
    "Internal": false,
    "Attachable": false,
    "Containers": {
      "156849ccd902b812b7d17f05d2d81532ccebe5bf788c9a79de63e12bb92fc621"
      : {
        "Name": "alpine4",
        "EndpointID":
        "7277c5183f0da5148b33d05f329371fce7befc5282d2619cfb23690b2adf467d"
        ,
        "MacAddress": "02:42:ac:11:00:03",
        "IPv4Address": "172.17.0.3/16",
        "IPv6Address": ""
      },
      "fa1340b8d83eef5497166951184ad3691eb48678a3664608ec448a687b047c53"
      : {
        "Name": "alpine3",
        "EndpointID":
        "5ae767367dcbebc712c02d49556285e888819d4da6b69d88cd1b0d52a83af95f"
        ,
        "MacAddress": "02:42:ac:11:00:02",
        "IPv4Address": "172.17.0.2/16",
        "IPv6Address": ""
      }
    },
    "Options": {
      "com.docker.network.bridge.default_bridge": "true",
      "com.docker.network.bridge.enable_icc": "true",
      "com.docker.network.bridge.enable_ip_masquerade":
      "true",
      "com.docker.network.bridge.host_binding_ipv4":
      "0.0.0.0",
      "com.docker.network.bridge.name": "docker0",
      "com.docker.network.driver.mtu": "1500"
    },
    "Labels": {}
  }
]

```

Containers **alpine3** and **alpine4** are connected to the bridge network.

```
$ docker network inspect alpine-net
```

```
[
  {
    "Name": "alpine-net",
    "Id":
    "e9261a8c9a19eabf2bf1488bf5f208b99b1608f330cff585c273d39481c9b0ec"
  ,
    "Created": "2017-09-25T21:38:12.620046142Z",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": {},
      "Config": [
        {
          "Subnet": "172.18.0.0/16",
          "Gateway": "172.18.0.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Containers": {

    "0a02c449a6e9a15113c51ab2681d72749548fb9f78fae4493e3b2e4e74199c4a"
    : {
      "Name": "alpine1",
      "EndpointID":
      "c83621678eff9628f4e2d52baf82c49f974c36c05cba152db4c131e8e7a64673"
    ,
      "MacAddress": "02:42:ac:12:00:02",
      "IPv4Address": "172.18.0.2/16",
      "IPv6Address": ""
    },

    "156849ccd902b812b7d17f05d2d81532ccebe5bf788c9a79de63e12bb92fc621"
    : {
      "Name": "alpine4",
      "EndpointID":
      "058bc6a5e9272b532ef9a6ea6d7f3db4c37527ae2625d1cd1421580fd0731954"
    ,
      "MacAddress": "02:42:ac:12:00:04",
      "IPv4Address": "172.18.0.4/16",
      "IPv6Address": ""
    },
  },

```

```
"a535d969081e003a149be8917631215616d9401edcb4d35d53f00e75ea1db653"
: {
    "Name": "alpine2",
    "EndpointID":
"198f3141ccf2e7dba67bce358d7b71a07c5488e3867d8b7ad55a4c695ebb8740"
,
    "MacAddress": "02:42:ac:12:00:03",
    "IPv4Address": "172.18.0.3/16",
    "IPv6Address": ""
}
},
"Options": {},
"Labels": {}
}
]
```

Containers **alpine1**, **alpine2**, and **alpine4** are connected to the alpine-net network.

On user-defined networks like alpine-net, containers can not only communicate by IP address, but can also resolve a container name to an IP address. This capability is called automatic service discovery. Let's connect to alpine1 and test this out. alpine1 should be able to resolve alpine2 and alpine4 (and alpine1, itself) to IP addresses.

```
$ docker container attach alpine1
```

```
# ping -c 2 alpine2
```

```
PING alpine2 (172.18.0.3): 56 data bytes
64 bytes from 172.18.0.3: seq=0 ttl=64 time=0.085 ms
64 bytes from 172.18.0.3: seq=1 ttl=64 time=0.090 ms
```

```
--- alpine2 ping statistics ---
```

```
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 0.085/0.087/0.090 ms
```

```
# ping -c 2 alpine4
```

```
PING alpine4 (172.18.0.4): 56 data bytes
64 bytes from 172.18.0.4: seq=0 ttl=64 time=0.076 ms
64 bytes from 172.18.0.4: seq=1 ttl=64 time=0.091 ms
```

```
--- alpine4 ping statistics ---
```

```
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 0.076/0.083/0.091 ms
```

```
# ping -c 2 alpine1
```

```
PING alpine1 (172.18.0.2): 56 data bytes
64 bytes from 172.18.0.2: seq=0 ttl=64 time=0.026 ms
64 bytes from 172.18.0.2: seq=1 ttl=64 time=0.054 ms
```

```
--- alpine1 ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 0.026/0.040/0.054 ms
```

From alpine1, you should not be able to connect to alpine3 at all, since it is not on the alpine-net network.

```
# ping -c 2 alpine3
```

```
ping: bad address 'alpine3'
```

Not only that, but you can't connect to alpine3 from alpine1 by its IP address either. Look back at the docker network inspect output for the bridge network and find alpine3's IP address: 172.17.0.2 Try to ping it.

```
# ping -c 2 172.17.0.2
```

```
PING 172.17.0.2 (172.17.0.2): 56 data bytes

--- 172.17.0.2 ping statistics ---
2 packets transmitted, 0 packets received, 100% packet loss
```

Detach from alpine1 using detach sequence, **CTRL + p CTRL + q (hold down CTRL and type p followed by q)**.

Remember that alpine4 is connected to both the default bridge network and alpine-net. It should be able to reach all of the other containers. However, you will need to address alpine3 by its IP address. Attach to it and run the tests.

```
$ docker container attach alpine4
```

```
# ping -c 2 alpine1
```

```
PING alpine1 (172.18.0.2): 56 data bytes
64 bytes from 172.18.0.2: seq=0 ttl=64 time=0.074 ms
64 bytes from 172.18.0.2: seq=1 ttl=64 time=0.082 ms
```

```
--- alpine1 ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 0.074/0.078/0.082 ms
```

```
# ping -c 2 alpine2
```

```
PING alpine2 (172.18.0.3): 56 data bytes  
64 bytes from 172.18.0.3: seq=0 ttl=64 time=0.075 ms  
64 bytes from 172.18.0.3: seq=1 ttl=64 time=0.080 ms
```

```
--- alpine2 ping statistics ---  
2 packets transmitted, 2 packets received, 0% packet loss  
round-trip min/avg/max = 0.075/0.077/0.080 ms
```

```
# ping -c 2 alpine3  
ping: bad address 'alpine3'
```

```
# ping -c 2 172.17.0.2
```

```
PING 172.17.0.2 (172.17.0.2): 56 data bytes  
64 bytes from 172.17.0.2: seq=0 ttl=64 time=0.089 ms  
64 bytes from 172.17.0.2: seq=1 ttl=64 time=0.075 ms
```

```
--- 172.17.0.2 ping statistics ---  
2 packets transmitted, 2 packets received, 0% packet loss  
round-trip min/avg/max = 0.075/0.082/0.089 ms
```

```
# ping -c 2 alpine4
```

```
PING alpine4 (172.18.0.4): 56 data bytes  
64 bytes from 172.18.0.4: seq=0 ttl=64 time=0.033 ms  
64 bytes from 172.18.0.4: seq=1 ttl=64 time=0.064 ms
```

```
--- alpine4 ping statistics ---  
2 packets transmitted, 2 packets received, 0% packet loss  
round-trip min/avg/max = 0.033/0.048/0.064 ms
```

As a final test, make sure your containers can all connect to the internet by pinging google.com. You are already attached to alpine4 so start by trying from there. Next, detach from alpine4 and connect to alpine3 (which is only attached to the bridge network) and try again. Finally, connect to alpine1 (which is only connected to the alpine-net network) and try again.

```
# ping -c 2 google.com
```

```
PING google.com (172.217.3.174): 56 data bytes  
64 bytes from 172.217.3.174: seq=0 ttl=41 time=9.778 ms  
64 bytes from 172.217.3.174: seq=1 ttl=41 time=9.634 ms
```



```
--- google.com ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 9.634/9.706/9.778 ms
```

```
CTRL+p CTRL+q
```

```
$ docker container attach alpine3
```

```
# ping -c 2 google.com
```

```
PING google.com (172.217.3.174): 56 data bytes
64 bytes from 172.217.3.174: seq=0 ttl=41 time=9.706 ms
64 bytes from 172.217.3.174: seq=1 ttl=41 time=9.851 ms
```

```
--- google.com ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 9.706/9.778/9.851 ms
```

```
CTRL+p CTRL+q
```

```
$ docker container attach alpine1
```

```
# ping -c 2 google.com
```

```
PING google.com (172.217.3.174): 56 data bytes
64 bytes from 172.217.3.174: seq=0 ttl=41 time=9.606 ms
64 bytes from 172.217.3.174: seq=1 ttl=41 time=9.603 ms
```

```
--- google.com ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
round-trip min/avg/max = 9.603/9.604/9.606 ms
```

```
CTRL+p CTRL+q
```

Stop and remove all containers and the alpine-net network.

```
$ docker container stop alpine1 alpine2 alpine3 alpine4
```

```
$ docker container rm alpine1 alpine2 alpine3 alpine4
```

```
$ docker network rm alpine-net
```

LAB-11

`docker commit [OPTIONS] CONTAINER [REPOSITORY[:TAG]]`

It can be useful to commit a container's file changes or settings into a new image. This allows you to debug a container by running an interactive shell, or to export a working dataset to another server. Generally, it is better to use **Dockerfiles** to manage your images in a documented and maintainable way

Start a new container based on fedora images to start a bash as entrypoint

```
[root@server ~]# docker run -dit fedora bash
```

```
465b93c17a36ec17cc54d053c41d27bb5822e1483d7facbb8c598f47e3385e66
```

```
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	
STATUS	PORTS	NAMES		
465b93c17a36	fedora	"bash"	5 seconds ago	Up 4
seconds		strange_khorana		

Commit your container

```
[root@server ~]# docker commit 465b93c17a36
```

```
class/testimage:version1
```

```
sha256:61b13e14d2377d6be7d850b7af1b0b4268045680149151f25313efa076d05050
```

```
root@server ~]# docker images
```

REPOSITORY	TAG	IMAGE ID
CREATED	SIZE	
class/testimage	version1	61b13e14d237
28 seconds ago	215MB	

Commit a container with new configuration

```
[root@server ~]# docker inspect -f "{{ .Config.Env }}" 465b93c17a36
```

```
[PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin  
]
```

```
[root@server ~]# docker commit --change "ENV AUTHOR=luca"  
465b93c17a36 class/testimage:version2  
sha256:f6e9d61b16cfaa10ece2afe1c73982edec0e5c1fcf9212014f0741ad1c6e50c
```

```
[root@server ~]# docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED
SIZE			
class/testimage	version2	f6e9d61b16cf	4 seconds ago
215MB			
class/testimage	version1	61b13e14d237	4 minutes ago
215MB			

```
[root@server ~]# docker inspect -f "{{ .Config.Env }}"  
f6e9d61b16cfaa10ece2afe1c73982edec0e5c1fcf9212014f0741ad1c6e50c
```

```
[PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin  
AUTHOR=luca]
```

Commit a container with new CMD and ENTRYPOINT instructions

```
[root@server ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
465b93c17a36	fedora	"bash"	12 minutes ago
Up 12 minutes		strange_khorana	

```
[root@server ~]# docker commit --change='ENTRYPOINT ["echo"]'  
--change='CMD ["hello", "world"]' 465b93c17a36  
class/testimage:version3  
sha256:2cc2d990114e577d92c662ed5a325a976732def93d5d31d0924cd032600  
6f82e
```

```
[root@server ~]# docker images
```

REPOSITORY	TAG	IMAGE ID
CREATED	SIZE	
class/testimage	version3	2cc2d990114e
7 seconds ago	215MB	

class/testimage	version2	f6e9d61b16cf
8 minutes ago	215MB	
class/testimage	version1	61b13e14d237

Run your version3 container

```
[root@server ~]# docker run class/testimage:version3
```

hello world

LAB-12

Create your “hello world” image

Creating a directory called `myhelloworld` where we'll create the following files:

- `start.sh`
- `Dockerfile`

Create the **start.sh** with the following content:

```
#!/bin/sh
echo "Hello $AUTHOR"
```

N.B. Make sure to add the execution permission to the script `start.sh`

Create the file **Dockerfile** based on `busybox` images

```
FROM busybox:latest
```

Copy the local `start.sh` into container under `/` directory

```
COPY start.sh /start.sh
```

Add the environment variable `AUTHOR` with default value `Docker`

```
ENV \
AUTHOR=Docker
```

The command for running the application will be now your `start.sh` script

```
CMD ["/start.sh"]
```

Your `Dockerfile` is now ready.

```
FROM busybox:latest
```

```
COPY start.sh /start.sh
```

```
ENV \
AUTHOR=Docker
```

```
CMD ["/start.sh"]
```

Build you image as follow:

```
docker build -t class/myhelloworld:v1 .
Sending build context to Docker daemon 3.072kB
```

```
Step 1/4 : FROM busybox:latest
---> 020584afccce
Step 2/4 : COPY start.sh /start.sh
---> Using cache
---> 6f9ffe05fe61
Step 3/4 : ENV AUTHOR=Docker
---> Using cache
---> 545f5ac64db5
Step 4/4 : CMD ["/start.sh"]
---> Using cache
---> 2989ecd8edd5
Successfully built 2989ecd8edd5
Successfully tagged class/myhelloworld:latest
```

Run you image!!

```
docker run -e AUTHOR="Mario Rossi" class/myhelloworld
```

```
Hello Mario Rossi
```

LAB-13

Docker Images

The goal of this lab is to create a Docker image which will run a [Flask](#) app. For the purposes of this exercise, we've a fun little Python Flask app that displays a random cat `.gif` every time it is loaded.

Start by creating a directory called `flask-app` where we'll create the following files:

- `app.py`
- `requirements.txt`
- `templates/index.html`
- `Dockerfile`

Make sure to `cd flask-app` before you start creating the files, because you don't want to start adding a whole bunch of other random files to your image.

app.py

Create the `app.py` with the following content:

```
from flask import Flask, render_template
import random

app = Flask(__name__)

# list of images
images = [

    "https://www.animatedimages.org/data/media/209/animated-cat-image-0072.gif",
    "https://www.animatedimages.org/data/media/209/animated-cat-image-0056.gif",
    "https://www.animatedimages.org/data/media/209/animated-cat-image-0394.gif",
    "https://www.animatedimages.org/data/media/209/animated-cat-image-0338.gif",
    "https://www.animatedimages.org/data/media/209/animated-cat-image-0058.gif",
    "https://www.animatedimages.org/data/media/209/animated-cat-image-0184.gif",
    "https://www.animatedimages.org/data/media/209/animated-cat-image-0459.gif"
]

@app.route('/')
def index():
    url = random.choice(images)
    return render_template('index.html', url=url)
```

```
if __name__ == "__main__":  
    app.run(host="0.0.0.0")
```

requirements.txt

In order to install the Python modules required for our app, we need to create a file called requirements.txt and add the following line to that file:

```
Flask==2.1.0
```

templates/index.html

Create a directory called templates and create an index.html file in that directory with the following content in it:

```
<html>  
  <head>  
    <style type="text/css">  
      body {  
        background: black;  
        color: white;  
      }  
      div.container {  
        max-width: 500px;  
        margin: 100px auto;  
        border: 20px solid white;  
        padding: 10px;  
        text-align: center;  
      }  
      h4 {  
        text-transform: uppercase;  
      }  
    </style>  
  </head>  
  <body>  
    <div class="container">  
      <h4>Cat Gif of the day</h4>  
        
      <br>  
    </div>  
  </body>
```



```
</html>
```

Write a Dockerfile

We want to create a Docker image with this web app. As mentioned above, all user images are based on a base image. Since our application is written in Python, we will build our own Python image based on [Alpine](#). We'll do that using a Dockerfile.

A [Dockerfile](#) is a text file that contains a list of commands that the Docker daemon calls while creating an image. The Dockerfile contains all the information that Docker needs to know to run the app — a base Docker image to run from, location of your project code, any dependencies it has, and what commands to run at start-up. It is a simple way to automate the image creation process. The best part is that the [commands](#) you write in a Dockerfile are almost identical to their equivalent Linux commands. This means you don't really have to learn new syntax to create your own Dockerfiles.

1. Create a file called Dockerfile, and add content to it as described below.

We'll start by specifying our base image, using the FROM keyword:

```
FROM alpine:latest
```

2. The next step usually is to write the commands of copying the files and installing the dependencies. But first we will install the Python pip package to the alpine linux distribution. This will not just install the pip package but any other dependencies too, which includes the python interpreter. Add the following [RUN](#) command next:

```
RUN apk add --update py3-pip
```

3. Let's add the files that make up the Flask Application.

Install all Python requirements for our app to run. This will be accomplished by adding the lines:

```
COPY requirements.txt /usr/src/app/  
RUN pip install --no-cache-dir -r /usr/src/app/requirements.txt
```

Copy the files you have created earlier into our image by using [COPY](#) command.

```
COPY app.py /usr/src/app/  
COPY templates/index.html /usr/src/app/templates/
```

4. Specify the port number which needs to be exposed. Since our flask app is running on 5000 that's what we'll expose.

```
EXPOSE 5000
```

5. The last step is the command for running the application which is simply - `python ./app.py`. Use the **CMD** command to do that:

```
CMD ["python", "/usr/src/app/app.py"]
```

The primary purpose of `CMD` is to tell the container which command it should run by default when it is started.

6. Verify your Dockerfile.

Our `Dockerfile` is now ready. This is how it looks:

```
FROM alpine:latest

# Install python and pip

RUN apk add --update py3-pip

# install Python modules needed by the Python app

COPY requirements.txt /usr/src/app/

#install all requirements.txt packages

RUN pip install --no-cache-dir -r /usr/src/app/requirements.txt

# copy files required for the app to run

COPY app.py /usr/src/app/

COPY templates/index.html /usr/src/app/templates/

# tell the port number the container should expose

EXPOSE 5000

# run the application

CMD ["python3", "/usr/src/app/app.py"]
```

Build the image

Now that you have your `Dockerfile`, you can build your image. The `docker build` command does the heavy-lifting of creating a docker image from a `Dockerfile`. The `docker build` command is quite simple - it takes an optional tag name with the `-t` flag, and the location of the directory containing the `Dockerfile` - the `.` indicates the current directory:

```
[root@server flask-app]# docker build -t class/myfirstapp .
Sending build context to Docker daemon 7.68kB
Step 1/8 : FROM alpine:3.5
3.5: Pulling from library/alpine
8cae0e1ac61c: Pull complete
Digest:
sha256:66952b313e51c3bd1987d7c4ddf5dba9bc0fb6e524eed2448fa6602
46b3e76ec
Status: Downloaded newer image for alpine:3.5
---> f80194ae2e0c
Step 2/8 : RUN apk add --update py2-pip
---> Running in 1fb2f28d3369
fetch
http://dl-cdn.alpinelinux.org/alpine/v3.5/main/x86_64/APKINDEX
.tar.gz
fetch
http://dl-cdn.alpinelinux.org/alpine/v3.5/community/x86_64/APK
INDEX.tar.gz
(1/12) Installing libbz2 (1.0.6-r5)
(2/12) Installing expat (2.2.0-r1)
[...]
```

Successfully built e78e604f9de2
Successfully tagged class/myfirstapp:latest

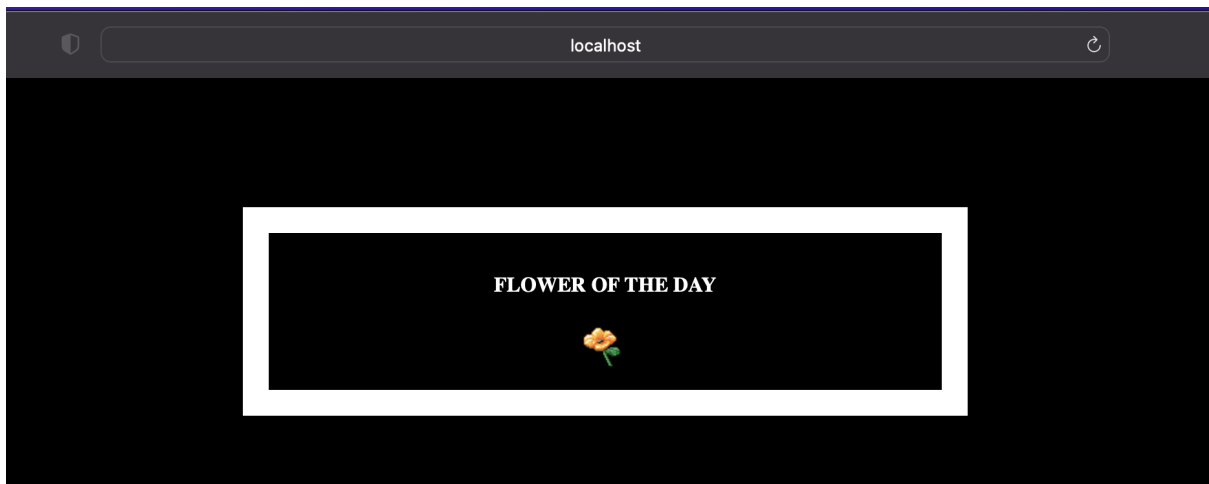
```
[root@server flask-app]# docker images
```

REPOSITORY	TAG	IMAGE ID
class/myfirstapp	latest	e78e604f9de2
11 seconds ago	57MB	

Run your image

```
[root@server flask-app]# docker run -p 80:5000 --name  
myfirstapp class/myfirstapp  
* Running on http://0.0.0.0:5000/ (Press CTRL+C to quit)
```

Open a http connection from your local browser on forwarded port 8080 (virtualbox system)



Refresh the page to get a new random gif image!!

LAB-14

How to quickly replace environment variables in a file

The [envsubst](#) is part of the [gettext](#) internationalization and localization project for unix.

Example

Let's say, we have an existing configuration file, that want to use to configure some application at the start of container.

```
# my configuration file
server: https://someurl.com/auth
username: foo_user
password: foo_password
```

1. Create sample configuration file

Lets replace the information with the environment variables:

```
server: $SERVER_URL
username: $USER_NAME
password: $USER_PASSWORD
```

2. Create the ENV file

```
SERVER_URL=https://someurl.com/auth
USER_NAME=foo_user
USER_PASSWORD=foo_password
```

3. Substitution

To run an actual substitution, perform the following commands.

```
envsubst < config.txt > config.conf
```

```
cat config.conf
server: https://someurl.com/auth
username: foo_user
password: foo_password
```

Create your own static docker httpd site based on fedora

Dockerfile:

```
FROM fedora
ENV AUTOR=Docker
RUN yum -y install httpd gettext
WORKDIR /var/www/html
COPY Hello_docker.html /var/www/html
EXPOSE 80
CMD cd /var/www/html && envsubst < Hello_docker.html > index.html;
/usr/sbin/httpd -D FOREGROUND;
```

Hello_docker.html:

```
<!DOCTYPE html>
<html>
<head>
<style>
body {
    margin-left: 500px;
    background: #5d9ab2
url("https://www.docker.com/sites/default/files/d8/styles/role_ico
n/public/2019-07/Docker-Logo-White-RGB_Vertical-BG_0.png?itok=8Tua
c9I3") no-repeat top left;
}

.center_div {
    border: 1px solid gray;
    margin-left: auto;
    margin-right: auto;
    width: 90%;
    background-color: #d0f0f6;
    text-align: left;
    padding: 8px;
}
</style>
</head>
<body>

<div class="center_div">
    <h1>Hello $AUTHOR</h1>
</div>
```

</body>

</html>

Build you image:

```
docker build -t dockersample/lab-14:1.0 .
```

Run a new container based on your image

```
docker run -d -e AUTHOR="luca" -p 80:80 dockersample/lab-14:1.0
```

check that your container is up and running

```
docker ps
```

Check on your local browser

<http://localhost:8080>

LAB-15

no solution for this labs

1)

Create a new images based on:

- Operating System “fedora”
- install httpd service
- expose your container to 80 port
- modify /etc/httpd/conf/httpd.conf to use /srv/www/http/ as DocumentRoot directory
- Add in Dockerfile a httpd start command at the end of file with content:
 - **CMD ["/usr/sbin/httpd","-D","FOREGROUND"]**
- Build the image
- Start your container bind a **local index.html** to /srv/www/http/index.html
- Try if httpd works from browser

2)

- Create a new fedora based Docker image called yourName/hello:1.0
- Once launched, the task of this image will be print the string "Hello Student" on the screen every 2 seconds, for a maximum of 5 times, and then exit by printing the string "goodbye !!"
- The word Student must be modified through the environment variable passed at the start of the container.

3)

- Modify **LAB-14**
 - remove from Dockerfile
 - `CMD cd /var/www/html && envsubst < Hello_docker.html > index.html ; /usr/sbin/httpd -D FOREGROUND;`
 - use a ENTRYPOINT to call a bash script called `entrypoint.sh`
 - the script will have to:
 - replace the environment variable inside the index.html.
 - start the httpd server

4)

- Modify point 1 to configure dynamically the httpd at container start:
- Create these two env
 - DOCUMENTROOT
 - PORT
- Add these variable into httpd.conf to configure the listen port and documentRoot entry (do not forget Directory configuration: `<Directory "<DocumentRoot directory>"`)

- **Run example:**

```
docker run -d -p 80:8888 --rm -e  
"DOCUMENTROOT=/srv/lab15-3/html" -e "PORT=8888" -v  
$PWD/index:/srv/lab15-3/html yourImageName:latest
```

LAB-16

Deploy a registry server

Use a command like the following to start the registry container:

```
$ docker run -d -p 5000:5000 --restart=always --name registry
registry:2
```

Copy an image from Docker Hub to your registry

You can pull an image from Docker Hub and push it to your registry. The following example pulls the `ubuntu:16.04` image from Docker Hub and re-tags it as `my-ubuntu`, then pushes it to the local registry. Finally, the `ubuntu:16.04` and `my-ubuntu` images are deleted locally and the `my-ubuntu` image is pulled from the local registry.

Pull the `ubuntu:16.04` image from Docker Hub.

```
$ docker pull ubuntu:16.04
```

Tag the image as `localhost:5000/my-ubuntu`. This creates an additional tag for the existing image. When the first part of the tag is a hostname and port, Docker interprets this as the location of a registry, when pushing.

```
$ docker tag ubuntu:16.04 localhost:5000/my-ubuntu
```

Push the image to the local registry running at `localhost:5000`:

```
$ docker push localhost:5000/my-ubuntu
```

Remove the locally-cached `ubuntu:16.04` and `localhost:5000/my-ubuntu` images, so that you can test pulling the image from your registry. This does not remove the `localhost:5000/my-ubuntu` image from your registry.

```
$ docker image remove ubuntu:16.04
$ docker image remove localhost:5000/my-ubuntu
```

Pull the `localhost:5000/my-ubuntu` image from your local registry.

```
$ docker pull localhost:5000/my-ubuntu
```

Customize the published port

If you are already using port 5000, or you want to run multiple local registries to separate areas of concern, you can customize the registry's port settings. This example runs the registry on port 5001 and also names it `registry-test`. Remember, the first part of the `-p` value is the host port and the second part is the port within the container. Within the container, the registry listens on port 5000 by default.

```
$ docker run -d \  
-p 5001:5000 \  
--name registry-test \  
registry:2
```

If you want to change the port the registry listens on within the container, you can use the environment variable `REGISTRY_HTTP_ADDR` to change it. This command causes the registry to listen on port 5001 within the container:

```
$ docker run -d \  
-e REGISTRY_HTTP_ADDR=0.0.0.0:5001 \  
-p 5001:5001 \  
--name registry-test \  
registry:2
```

Storage customization

Customize the storage location

By default, your registry data is persisted as a [docker volume](#) on the host filesystem. If you want to store your registry contents at a specific location on your host filesystem, such as if you have an SSD or SAN mounted into a particular directory, you might decide to use a bind mount instead. A bind mount is more dependent on the filesystem layout of the Docker host, but more performant in many situations. The following example bind-mounts the host directory `/mnt/registry` into the registry container at `/var/lib/registry/`.

```
$ docker run -d \  
-p 5000:5000 \  
--restart=always \  
--name registry \  
-v /mnt/registry:/var/lib/registry \  
registry:2
```

COMPOSE

LAB-17

Run and build your hello world docker images created on LAB-12 with docker-compose. Make sure you have already installed both [Docker Engine](#) and [Docker Compose](#). Create a directory called lab17

Create a directory called lab17/myhelloworld with these files:

- start.sh
- Dockerfile

Create the **start.sh** with the following content:

```
#!/bin/sh
echo "Hello $AUTHOR"
```

N.B. Make sure to add the execution permission to the script start.sh

Create the file **Dockerfile** based on `busybox` images

```
FROM busybox:latest
COPY start.sh /start.sh
RUN chmod +x /start.sh
ENV AUTHOR=Docker
CMD ["/start.sh"]
```

Create the **docker-compose.yml** file under lab17 directory

```
version: "3.8"
services:
  hello:
    build: ./myhelloworld
    image: "compose/hello-world"
    environment:
      - AUTHOR=Docker
```

start it with docker compose tool:

```
docker compose up -d
```

LAB-18

create your own docker-compose.yml file to manage **LAB-14**

LAB-19

Compose and WordPress

Create an empty project directory.

You can name the directory something easy for you to remember. This directory is the context for your application image. The directory should only contain resources to build that image.

This project directory contains a **docker-compose.yml** file

For example, if you named your directory my_wordpress:

```
$ cd my_wordpress/
```

Create a **docker-compose.yml** file that starts your WordPress blog and a separate MySQL instance with volume mounts for data persistence:

```
services:
  db:
    image: mariadb:10.6.4-focal
    command:
      '--default-authentication-plugin=mysql_native_password'
    volumes:
      - db_data:/var/lib/mysql
    restart: always
    environment:
      - MYSQL_ROOT_PASSWORD=somewordpress
      - MYSQL_DATABASE=wordpress
      - MYSQL_USER=wordpress
      - MYSQL_PASSWORD=wordpress
    expose:
      - 3306
      - 33060
  wordpress:
    image: wordpress:latest
    ports:
      - 80:80
    restart: always
    environment:
      - WORDPRESS_DB_HOST=db
      - WORDPRESS_DB_USER=wordpress
      - WORDPRESS_DB_PASSWORD=wordpress
```



```
    - WORDPRESS_DB_NAME=wordpress
volumes:
  db_data:
```

Build the project

Now, run **docker compose up -d** from your project directory.

This runs docker compose up in detached mode, pulls the needed Docker images, and starts the wordpress and database containers, as shown in the example below.

```
$ docker compose up -d
```

Creating network "my_wordpress_default" with the default driver

Pulling db (mysql:5.7)...

5.7: Pulling from library/mysql

efd26ecc9548: Pull complete

a3ed95caeb02: Pull complete

<...>

Bring up WordPress in a web browser

<http://localhost:8080>

At this point, WordPress should be running on port 80 of your Docker Host, and you can complete the "famous five-minute installation" as a WordPress administrator.



English (United States)

العربية المغربية

العربية

Azərbaycan dili

گۆنئی آذربایجان

Български

বাংলা

Bosanski

Català

Cebuano

Cymraeg

Dansk

Deutsch (Schweiz)

Deutsch

Continue



Welcome

Welcome to the famous five-minute WordPress installation process! Just fill in the information below and you'll be on your way to using the most extendable and powerful personal publishing platform in the world.

Information needed

Please provide the following information. Don't worry, you can always change these settings later.

Site Title

Username

Username can have only alphanumeric characters, spaces, underscores, hyphens, periods, and the @ symbol.

Password

Strong

Show

Important: You will need this password to log in. Please store it in a secure location.

Your Email

Double-check your email address before continuing.

Search Engine
Visibility

☐ Discourage search engines from indexing this site

It is up to search engines to honor this request.

Install WordPress

Shutdown and cleanup

The command `docker-compose down` removes the containers and default network, but preserves your WordPress database.

The command `docker-compose down --volumes` removes the containers, default network, and the WordPress database.

LAB-20

On this page you build a simple Python web application running on Docker Compose. The application uses the Flask framework and maintains a hit counter in Redis. While the sample uses Python, the concepts demonstrated here should be understandable even if you're not familiar with it.

Prerequisites

You don't need to install Python or Redis, as both are provided by Docker images.

Step 1: Setup

Define the application dependencies.

Create a directory for the project:

```
$ mkdir composetest
$ cd composetest
```

Create a file called **app.py** in your project directory and paste this in:

```
import time

import redis
from flask import Flask

app = Flask(__name__)
cache = redis.Redis(host='redis', port=6379)

def get_hit_count():
    retries = 5
    while True:
        try:
            return cache.incr('hits')
        except redis.exceptions.ConnectionError as exc:
            if retries == 0:
                raise exc
            retries -= 1
            time.sleep(0.5)

@app.route('/')
def hello():
    count = get_hit_count()
    return 'Hello World! I have been seen {}
times.\n'.format(count)
```

In this example, redis is the hostname of the redis container on the application's network. We

use the default port for Redis, 6379.

Handling transient errors

Note the way the `get_hit_count` function is written. This basic retry loop lets us attempt our request multiple times if the redis service is not available. This is useful at startup while the application comes online, but also makes our application more resilient if the Redis service needs to be restarted anytime during the app's lifetime. In a cluster, this also helps handling momentary connection drops between nodes.

Create another file called **requirements.txt** in your project directory and paste this in:

```
flask
redis
```

Step 2: Create a Dockerfile

In this step, you write a Dockerfile that builds a Docker image. The image contains all the dependencies the Python application requires, including Python itself.

In your project directory, create a file named **Dockerfile** and paste the following:

```
FROM python:3.7-alpine
WORKDIR /code
ENV FLASK_APP=app.py
ENV FLASK_RUN_HOST=0.0.0.0
RUN apk add --no-cache gcc musl-dev linux-headers
COPY requirements.txt requirements.txt
RUN pip install -r requirements.txt
EXPOSE 5000
COPY . .
CMD ["flask", "run"]
```

This tells Docker to:

- Build an image starting with the Python 3.7 image.
- Set the working directory to `/code`.
- Set environment variables used by the flask command.
- Install gcc and other dependencies
- Copy `requirements.txt` and install the Python dependencies.
- Add metadata to the image to describe that the container is listening on port 5000
- Copy the current directory `.` in the project to the `workdir` `.` in the image.
- Set the default command for the container to `flask run`.

Step 3: Define services in a Compose file

Create a file called **docker-compose.yml** in your project directory and paste the following:

```
version: "3.8"
services:
  web:
    build: .
    ports:
```

```
    - "80:5000"
redis:
  image: "redis:alpine"
```

This Compose file defines two services: web and redis.

Web service

The web service uses an image that's built from the Dockerfile in the current directory. It then binds the container and the host machine to the exposed port, 5000. This example service uses the default port for the Flask web server, 5000.

Redis service

The redis service uses a public [Redis](#) image pulled from the Docker Hub registry.

Step 4: Build and run your app with Compose

From your project directory, start up your application by running **docker-compose up**.

```
$ docker compose up
```

```
Creating network "composetest_default" with the default driver
Creating composetest_web_1 ...
Creating composetest_redis_1 ...
Creating composetest_web_1
Creating composetest_redis_1 ... done
Attaching to composetest_web_1, composetest_redis_1
web_1      | * Running on http://0.0.0.0:5000/ (Press CTRL+C to quit)
redis_1    | 1:C 17 Aug 22:11:10.480 # o000o000o000o Redis is starting
o000o000o000o
redis_1    | 1:C 17 Aug 22:11:10.480 # Redis version=4.0.1, bits=64,
commit=00000000, modified=0, pid=1, just started
redis_1    | 1:C 17 Aug 22:11:10.480 # Warning: no config file
specified, using the default config. In order to specify a config
file use redis-server /path/to/redis.conf
web_1      | * Restarting with stat
redis_1    | 1:M 17 Aug 22:11:10.483 * Running mode=standalone,
port=6379.
redis_1    | 1:M 17 Aug 22:11:10.483 # WARNING: The TCP backlog
setting of 511 cannot be enforced because
/proc/sys/net/core/somaxconn is set to the lower value of 128.
web_1      | * Debugger is active!
redis_1    | 1:M 17 Aug 22:11:10.483 # Server initialized
redis_1    | 1:M 17 Aug 22:11:10.483 # WARNING you have Transparent
Huge Pages (THP) support enabled in your kernel. This will create
latency and memory usage issues with Redis. To fix this issue run
the command 'echo never >
/sys/kernel/mm/transparent_hugepage/enabled' as root, and add it to
```


```
your /etc/rc.local in order to retain the setting after a reboot.  
Redis must be restarted after THP is disabled.  
web_1      | * Debugger PIN: 330-787-903  
redis_1    | 1:M 17 Aug 22:11:10.483 * Ready to accept connections
```

Compose pulls a Redis image, builds an image for your code, and starts the services you defined. In this case, the code is statically copied into the image at build time.

Enter `http://localhost:8080/` in a browser to see the application running.

You should see a message in your browser saying:


Hello World! I have been seen 1 times.



Refresh the page.

The number should increment.

Hello World! I have been seen 2 times.



Stop the application, either by running `docker-compose down` from within your project directory in the second terminal, or by hitting **CTRL+C** in the original terminal where you started the app.

Step 5: Edit the Compose file to add a bind mount

Edit docker-compose.yml in your project directory to add a [bind mount](#) for the web service:

```
version: "3.8"
services:
  web:
    build: .
    ports:
      - "80:5000"
    volumes:
      - ./code
    environment:
      FLASK_ENV: development
  redis:
    image: "redis:alpine"
```

The new volumes key mounts the project directory (current directory) on the host to /code inside the container, allowing you to modify the code on the fly, without having to rebuild the image. The environment key sets the FLASK_ENV environment variable, which tells flask run to run in development mode and reload the code on change. This mode should only be used in development.

Step 6: Re-build and run the app with Compose

From your project directory, type docker-compose up to build the app with the updated Compose file, and run it.

```
$ docker compose up
```

```
Creating network "composetest_default" with the default driver
Creating composetest_web_1 ...
Creating composetest_redis_1 ...
Creating composetest_web_1
Creating composetest_redis_1 ... done
Attaching to composetest_web_1, composetest_redis_1
web_1      | * Running on http://0.0.0.0:5000/ (Press CTRL+C to
quit)
...
```


Check the Hello World message in a web browser again, and refresh to see the count increment.

Because the application code is now mounted into the container using a volume, you can make changes to its code and see the changes instantly, without having to rebuild the image.

Change the greeting in **app.py** and save it. For example, change the Hello World! message to Hello from Docker!:

```
return 'Hello from Docker! I have been seen {}
times.\n'.format(count)
```


Refresh the app in your browser. The greeting should be updated, and the counter should still be incrementing.



Stop the application hitting **CTRL+C** in the original terminal where you started the app.

Step 8: Experiment with some other commands

If you want to run your services in the **background**, you can pass the **-d** flag (for “**detached**” mode) to `docker-compose up` and use `docker-compose ps` to see what is currently running:

```
$ docker compose up -d
```

Starting composetest_redis_1...

Starting composetest_web_1...

```
$ docker compose ps
```

Name	Command	State	Ports

-			
composetest_redis_1	/usr/local/bin/run	Up	
composetest_web_1	/bin/sh -c python app.py	Up	80>5000/tcp

The `docker-compose run` command allows you to run one-off commands for your services. For example, to see what environment variables are available to the web service:

```
$ docker-compose run web env
```

If you started Compose with `docker-compose up -d`, stop your services once you’ve finished with them:

```
$ docker-compose stop
```

You can bring everything down, removing the containers entirely, with the `down` command. Pass `--volumes` to also remove the data volume used by the Redis container:

```
$ docker compose down --volumes
```


LAB-21

try some example of docker compose

<https://github.com/docker/awesome-compose>

